

# ROLAND

Framework for extending static GNNs to dynamic

Steven Morse

William & Mary, Data Science  
DATA 691, Graph Learning


April 10, 2025

# Overview

*“ROLAND: Graph Learning Framework for Dynamic Graphs.” Jiaxuan You, Tianyu Du, Jure Leskovec (Stanford), KDD 2022.*

## Idea

How can we adapt powerful GNN architectures **directly** to a dynamic setting?



# Introduction and Related work

Adapting GNNs to dynamic setting introduces 3 challenges:

1. **Model design.** Sequence model on top of a GNN, RNN with GCN elements, ...
2. **Evaluation setting.** Partitioning temporal data risks distribution shifts.
3. **Training strategy.** Prohibitive memory overhead for dynamic training.


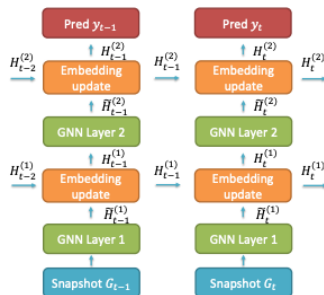



Figure: EvolveGCN

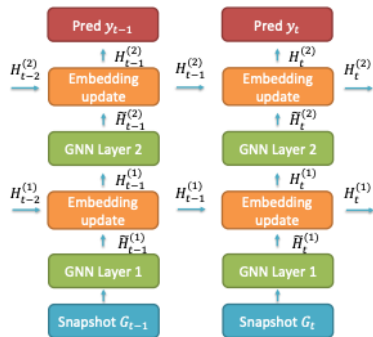
# Extending static to dynamic out-of-the-box

## Idea

Instead of doing dynamic updates at the top-level, do them **hierarchically** at each layer.



# A little more detail




---

## Algorithm 1 ROLAND GNN forward computation

---

**Input:** Dynamic graph snapshot  $G_t$ , hierarchical node state  $H_{t-1}$

**Output:** Prediction  $y_t$ , updated node state  $H_t$

---

- 1:  $H_t^{(0)} \leftarrow X_t$  {Initialize embedding from  $G_t$ }
  - 2: **for**  $l = 1, \dots, L$  **do**
  - 3:    $\tilde{H}_t^{(l)} = \text{GNN}^{(l)}(H_t^{(l-1)})$  {Implemented as Equation (4)}
  - 4:    $H_t^{(l)} = \text{UPDATE}^{(l)}(H_{t-1}^{(l)}, \tilde{H}_t^{(l)})$  {Equation (2)}
  - 5:  $y_t = \text{MLP}(\text{CONCAT}(\mathbf{h}_{u,t}^{(L)}, \mathbf{h}_{v,t}^{(L)})), \forall (u, v) \in E$  {Equation (5)}
- 

“GNN” is:

$$\mathbf{m}_{u \rightarrow v}^{(l)} = \mathbf{W}^{(l)} \text{CONCAT}^{(l)}(\mathbf{h}_u^{(l-1)}, \mathbf{h}_v^{(l-1)}, \mathbf{f}_{uv})$$

$$\mathbf{h}_v^{(l)} = \text{AGG}^{(l)}(\{\mathbf{m}_{u \rightarrow v}^{(l)} | u \in N(v)\}) + \mathbf{h}_v^{(l-1)}$$

# How to update?

- Moving average:

$$H_{t,v} = \kappa_{t,v} H_{t-1,v}^{(l)} + (1 - \kappa_{t,v}) \tilde{H}_{t,v}^{(l)}.$$

$$\kappa_{t,v} = \frac{\sum_{\tau=1}^{t-1} |E_{\tau}|}{\sum_{\tau=1}^{t-1} |E_{\tau}| + |E_t|}$$

- MLP:  $H_t^{(l)} = \text{MLP}(\text{CONCAT}(H_{t-1}^{(l)}, \tilde{H}_t^{(l-1)}))$
- GRU:  $H_t^{(l)} = \text{GRU}(H_{t-1}^{(l)}, \tilde{H}_t^{(l-1)})$




Figure: GRU

## Wait — isn't this the same as EvolveGCN?

*“While many works combine GNNs with recurrent models, we want to emphasize that our ROLAND framework is quite different.”*

**EvolveGCN:** weight params evolve (via RNN)

$$\begin{aligned}H_t^{(l+1)} &= \text{GNN}(H_t^{(l)}; W_t^{(l)}) \\ W_t^{(l)} &= \text{UPDATE}(H_t^{(l)}, W_{t-1}^{(l)})\end{aligned}$$

Memory efficient, but historical information slowly decays and we're limited to simpler GNN architectures.

**ROLAND:** weight params are static (w/in  $t$ )

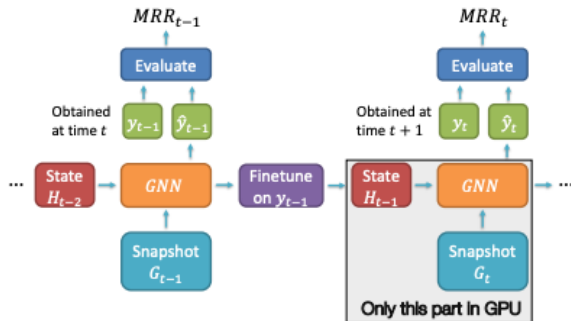
$$\begin{aligned}\tilde{H}_t^{(l+1)} &= \text{GNN}(H_t^{(l)}; W_t^{(l)}) \\ H_t^{(l)} &= \text{UPDATE}(H_{t-1}^{(l)}, \tilde{H}_t^{(l)})\end{aligned}$$

Not memory efficient (but fixed through their training strat), no decay, and no constraint on architecture.

# Live-update evaluation and meta-training

## Idea

Account for temporal distribution shifts by 1) training on the fly and 2) training a meta-model.



## Algorithm 2 ROLAND live-update evaluation

**Input:** Dynamic graph  $\mathcal{G} = \{G_1, \dots, G_T\}$ , link prediction labels  $y_1, \dots, y_T$ , number of snapshots  $T$ ,  $GNN(\cdot)$  defined in Algorithm 1  
**Output:** Performance MRR, model GNN

- 1: Initialize hierarchical node state  $H_0$
- 2: **for**  $t = 2, \dots, T$  **do**
- 3:   Collect link prediction labels  $y_{t-1} = y_{t-1}^{(train)} \cup y_{t-1}^{(val)}, y_t$
- 4:   **while**  $MRR_{t-1}^{(val)}$  is increasing **do**
- 5:      $H_{t-1}, \hat{y}_{t-1} \leftarrow GNN(G_{t-1}, H_{t-2}), \hat{y}_{t-1} = \hat{y}_{t-1}^{(train)} \cup \hat{y}_{t-1}^{(val)}$
- 6:     Update GNN via backprop based on  $\hat{y}_{t-1}^{(train)}, y_{t-1}^{(train)}$
- 7:      $MRR_{t-1}^{(val)} \leftarrow \text{EVALUATE}(\hat{y}_{t-1}^{(val)}, y_{t-1}^{(val)})$
- 8:      $H_t, \hat{y}_t \leftarrow GNN(G_t, H_{t-1})$
- 9:      $MRR_t \leftarrow \text{EVALUATE}(\hat{y}_t, y_t)$
- 10:  $MRR = \sum_{t=2}^T MRR_t / (T - 1)$

## Algorithm 3 ROLAND training algorithm

**Input:** Graph snapshot  $G_t$ , link prediction label  $y_t$ , hierarchical node state  $H_{t-1}$ , smoothing factor  $\alpha$ , meta-model  $GNN^{(meta)}$   
**Output:** Model GNN, updated meta-model  $GNN^{(meta)}$

- 1:  $GNN \leftarrow GNN^{(meta)}$
- 2: Move GNN,  $G_t, H_{t-1}$  to GPU
- 3: **while**  $MRR_t^{(val)}$  is increasing **do**
- 4:    $H_t, \hat{y}_t \leftarrow GNN(G_t, H_{t-1}), \hat{y}_t = \hat{y}_t^{(train)} \cup \hat{y}_t^{(val)}$
- 5:   Update GNN via backprop based on  $\hat{y}_t^{(train)}, y_t^{(train)}$
- 6:    $MRR_t^{(val)} \leftarrow \text{EVALUATE}(\hat{y}_t^{(val)}, y_t^{(val)})$
- 7: Remove GNN,  $G_t, H_{t-1}$  from GPU
- 8:  $GNN^{(meta)} \leftarrow (1 - \alpha)GNN^{(meta)} + \alpha GNN$



# Experiments

	Bitcoin-OTC	Bitcoin-Alpha	UCI-Message
GCN [14]	0.0025	0.0031	0.1141
DynGEM [12]	0.0921	0.1287	0.1055
dyngraph2vecAE [8]	0.0916	0.1478	0.0540
dyngraph2vecAERNN [8]	<b>0.1268</b>	<b>0.1945</b>	0.0713
EvolveGCN-H [30]	0.0690	0.1104	0.0899
EvolveGCN-O [30]	0.0968	0.1185	<b>0.1379</b>
ROLAND Moving Average	0.0468 $\pm$ 0.0022	0.1399 $\pm$ 0.0107	0.0649 $\pm$ 0.0049
ROLAND MLP	0.0778 $\pm$ 0.0024	0.1561 $\pm$ 0.0114	0.0875 $\pm$ 0.0110
ROLAND GRU	<b>0.2203 <math>\pm</math> 0.0167</b>	<b>0.2885 <math>\pm</math> 0.0123</b>	<b>0.2289 <math>\pm</math> 0.0618</b>
Improvement over best baseline	73.74%	43.33%	65.99%

Figure: Standard fixed-split setting

	BSI-ZK	AS-733	Reddit-Title	Reddit-Body	BSI-SVT	UCI-Message	Bitcoin-OTC	Bitcoin-Alpha
	MRR	MRR	MRR	MRR	MRR	MRR	MRR	MRR
Baseline Models with standard training								
EvolveGCN-H	N/A, OOM	N/A, OOM	N/A, OOM	<b>0.148 <math>\pm</math> 0.013</b>	<b>0.031 <math>\pm</math> 0.016</b>	0.061 $\pm$ 0.040	0.067 $\pm$ 0.035	0.079 $\pm$ 0.032
EvolveGCN-O	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	0.015 $\pm$ 0.006	0.071 $\pm$ 0.009	0.085 $\pm$ 0.022	0.071 $\pm$ 0.025
GCRN-GRU	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	0.080 $\pm$ 0.012	N/A, OOM	N/A, OOM
GCRN-LSTM	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	<b>0.083 <math>\pm</math> 0.001</b>	N/A, OOM	N/A, OOM
GCRN-Baseline	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	0.069 $\pm$ 0.004	<b>0.152 <math>\pm</math> 0.011</b>	<b>0.141 <math>\pm</math> 0.005</b>
TGCN	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	N/A, OOM	0.054 $\pm$ 0.024	0.128 $\pm$ 0.049	0.088 $\pm$ 0.038
Baseline Models with ROLAND Training								
EvolveGCN-H	N/A, OOM	0.251 $\pm$ 0.079	0.165 $\pm$ 0.026	0.102 $\pm$ 0.010	0.032 $\pm$ 0.008	0.057 $\pm$ 0.012	0.076 $\pm$ 0.022	0.054 $\pm$ 0.015
EvolveGCN-O	0.396	0.163 $\pm$ 0.002	0.047 $\pm$ 0.004	0.033 $\pm$ 0.001	0.018 $\pm$ 0.003	0.066 $\pm$ 0.012	0.032 $\pm$ 0.004	0.034 $\pm$ 0.002
GCRN-GRU	N/A, OOM	<b>0.344 <math>\pm</math> 0.001</b>	0.338 $\pm$ 0.006	0.217 $\pm$ 0.004	0.050 $\pm$ 0.004	0.089 $\pm$ 0.004	0.173 $\pm$ 0.003	0.140 $\pm$ 0.004
GCRN-LSTM	N/A, OOM	0.341 $\pm$ 0.001	0.344 $\pm$ 0.005	0.216 $\pm$ 0.000	0.051 $\pm$ 0.002	0.091 $\pm$ 0.010	0.174 $\pm$ 0.004	<b>0.146 <math>\pm</math> 0.005</b>
GCRN-Baseline	0.754	0.336 $\pm$ 0.002	0.351 $\pm$ 0.001	0.218 $\pm$ 0.002	0.054 $\pm$ 0.002	<b>0.095 <math>\pm</math> 0.008</b>	<b>0.183 <math>\pm</math> 0.002</b>	0.145 $\pm$ 0.003
TGCN	<b>0.831</b>	0.343 $\pm$ 0.002	<b>0.391 <math>\pm</math> 0.004</b>	<b>0.251 <math>\pm</math> 0.001</b>	<b>0.157 <math>\pm</math> 0.004</b>	0.080 $\pm$ 0.015	0.083 $\pm$ 0.011	0.069 $\pm$ 0.013
ROLAND results								
Moving Average	0.819	0.309 $\pm$ 0.011	0.362 $\pm$ 0.007	0.289 $\pm$ 0.038	0.177 $\pm$ 0.006	0.075 $\pm$ 0.006	0.120 $\pm$ 0.002	0.0962 $\pm$ 0.010
MLP-Update	0.834	0.329 $\pm$ 0.021	0.395 $\pm$ 0.006	0.291 $\pm$ 0.008	<b>0.217 <math>\pm</math> 0.003</b>	0.103 $\pm$ 0.010	0.154 $\pm$ 0.010	0.148 $\pm$ 0.012
GRU-Update	<b>0.851</b>	<b>0.340 <math>\pm</math> 0.001</b>	<b>0.425 <math>\pm</math> 0.015</b>	<b>0.362 <math>\pm</math> 0.002</b>	0.205 $\pm$ 0.014	<b>0.112 <math>\pm</math> 0.008</b>	<b>0.194 <math>\pm</math> 0.004</b>	<b>0.157 <math>\pm</math> 0.007</b>
Improvement over the best baseline	2.40%	-1.16%	8.70%	44.22%	38.21%	17.89%	6.01%	7.53%

Figure: Live update evaluation

# Experiments




Figure: ROLAND model retraining

# Subsequent and related work

- (Related) GraphGym framework (SNAP-Stanford)
- ROLAND = Discrete time, integrated. (See taxonomy)
- Extensions: WinGNN (2023)
- Applications: network biology (2024)




Figure: Dynamic graph learning taxonomy (from Zheng et al, 2025)




Figure: WinGNN (2023)

# Questions?